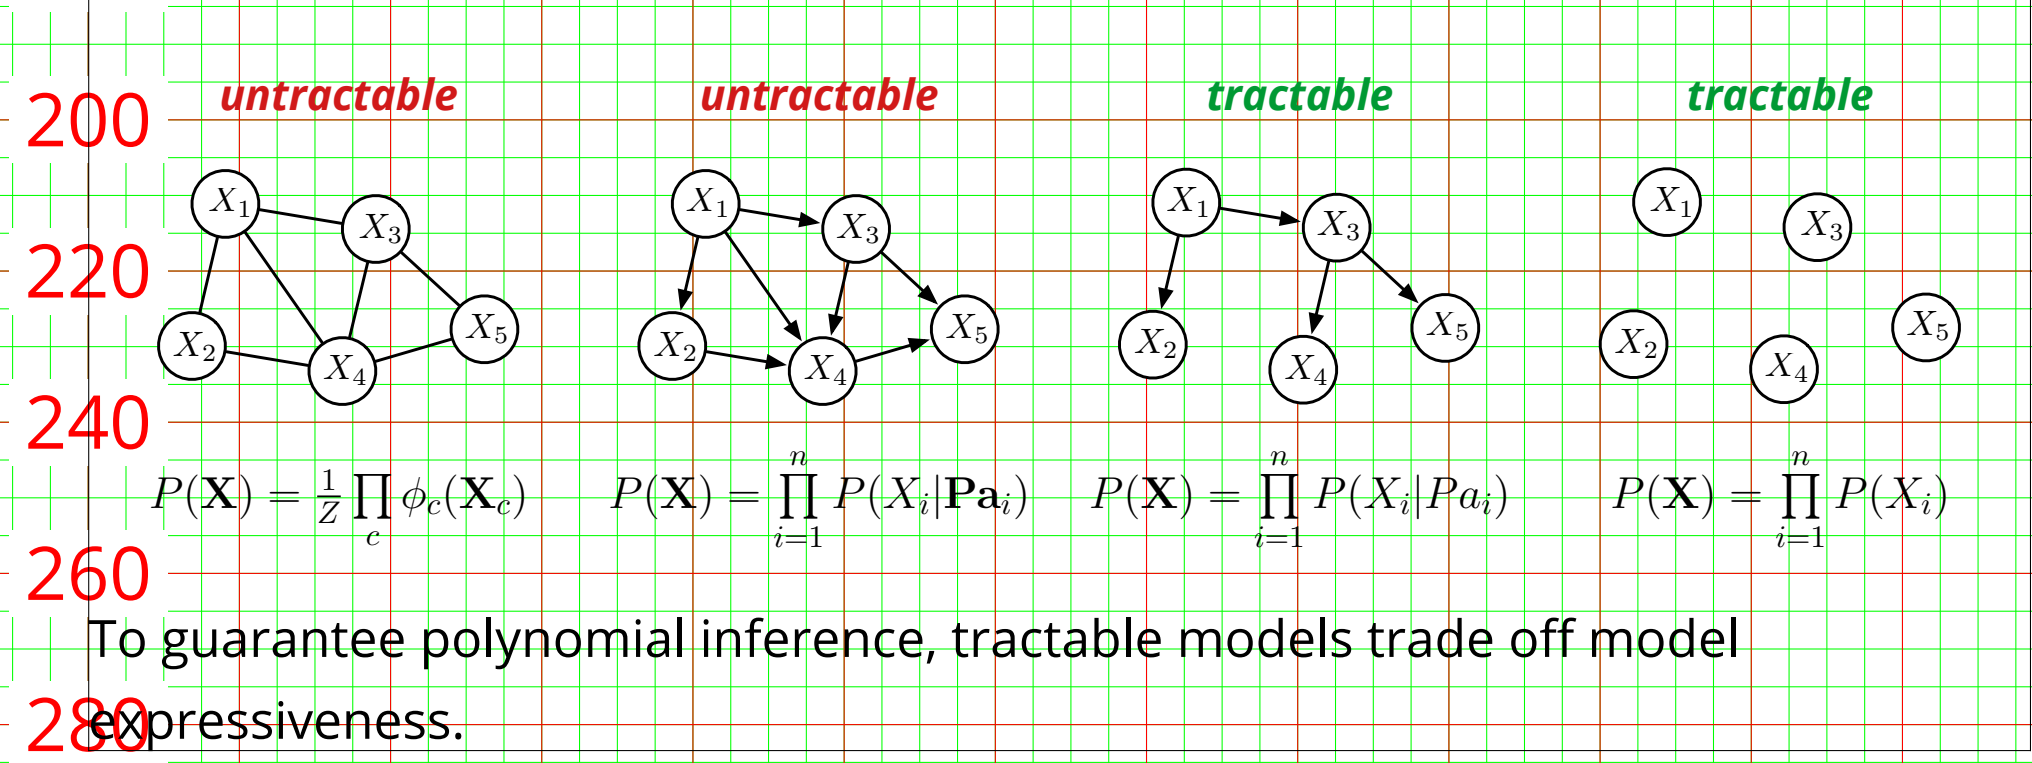


Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning

Antonio Vergari, Nicola Di Mauro and Floriana Esposito {firstname.lastname@uniba.it}

Sum-Product Networks and Tractable Models

Probabilistic Graphical Models (PGMs) provide a tool to compactly represent joint probability distributions $P(\mathbf{X})$. However, **inference**, the main task one may want to perform on a PGM, is generally **untractable**.




To guarantee polynomial inference, tractable models trade off model expressiveness.

Sum-Product Networks (SPNs) are DAGs *compiling* a pdf $P(\mathbf{X})$ into a **deep** architecture of **sum** and **product** nodes over univariate distributions X_1, \dots, X_n as leaves. The parameters of the network are the weights w_{ij} associated to sum nodes children edges.

Product nodes define factorizations over independent vars, sum nodes mixtures.

Products over nodes with different scopes (*decomposability*) and sums over nodes with same scopes (*completeness*) guarantee modeling a pdf (*validity*).



Bottom-up evaluation of the network:

$$S_{X_i}(x_j) = P(X_i = x_j)$$
$$S_{\oplus}(\mathbf{x}) = \sum_{i \in \text{ch}(\oplus)} w_i S_i(\mathbf{x}) \quad S_{\times}(\mathbf{x}) = \prod_{i \in \text{ch}(\times)} S_i(\mathbf{x})$$

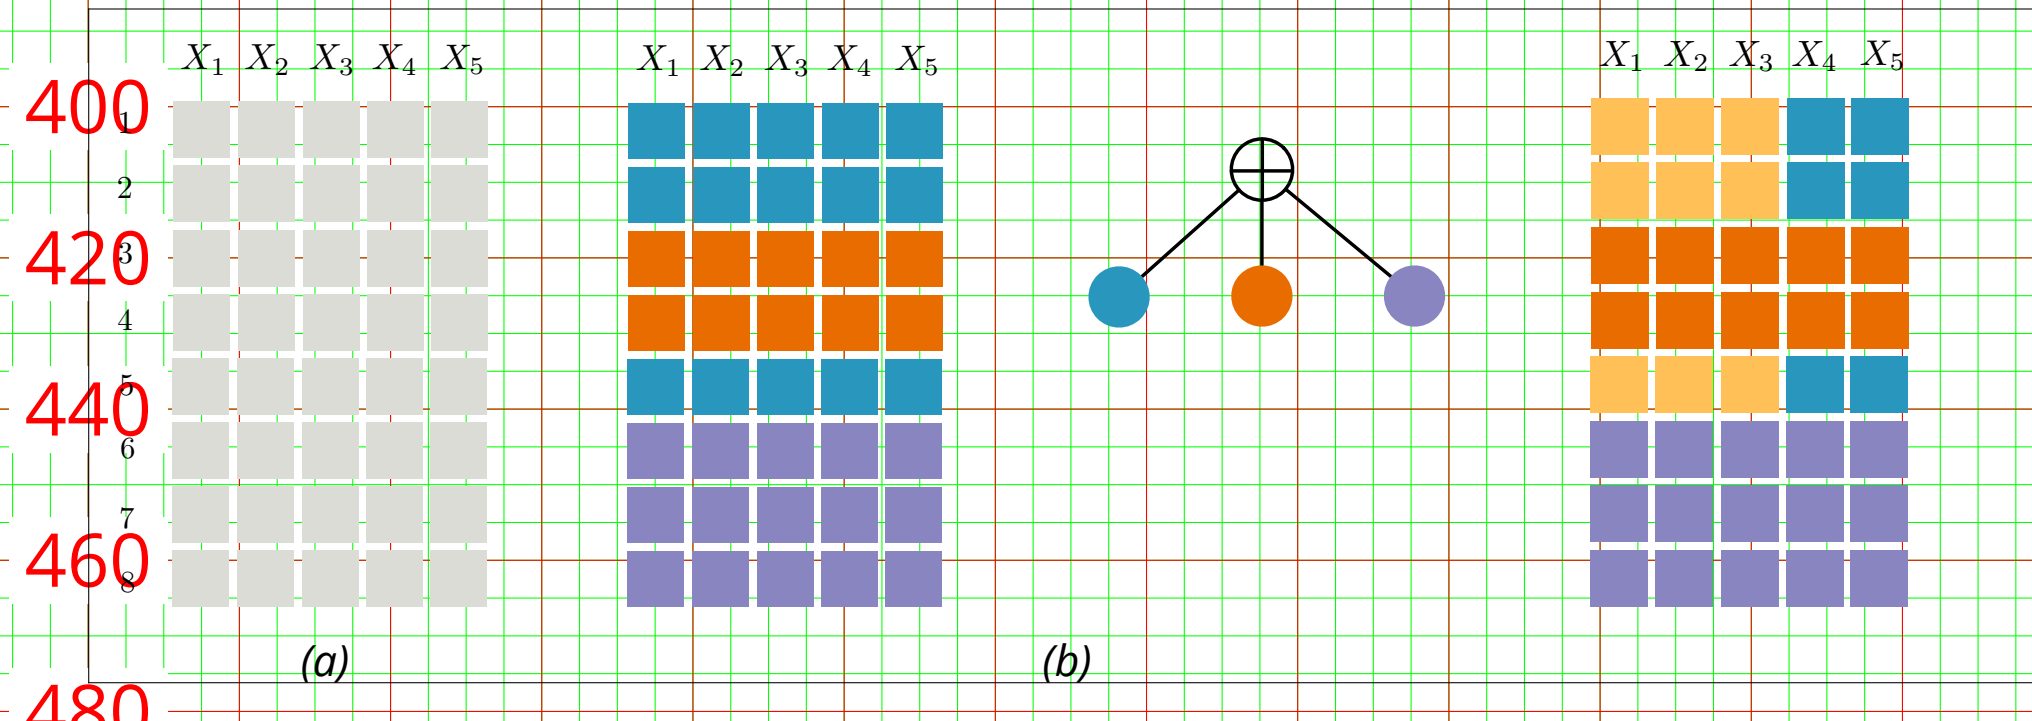
Inferences linear in the **size of the network** (# edges):

- $\oplus Z = S(\ast)$ (all leaves output 1)
- $\oplus P(\mathbf{e}) = S(\mathbf{e})/S(\ast)$
- $\oplus P(\mathbf{q}|\mathbf{e}) = \frac{P(\mathbf{q}, \mathbf{e})}{P(\mathbf{e})} = \frac{S(\mathbf{q}, \mathbf{e})}{S(\mathbf{e})}$
- $\oplus MPE(\mathbf{q}, \mathbf{e}) = \max_{\mathbf{q}} P(\mathbf{q}, \mathbf{e}) = S^{\max}(\mathbf{e})$, turning sum nodes into max nodes


The **depth of the network** (# layers) determines expressive efficiency [5, 9]

How and why to perform structure learning

SPN structure learning is a constraint-based search. Main ideas: to discover hidden variables for sum nodes and independences for product nodes by applying some form of clustering along matrix axis. Different variations: using K-Means on features [1]; merging features bottom-up with IB heuristics [7]; **LearnSPN** [2] is the first principled top-down greedy algorithm.



LearnSPN builds a tree-like SPN by recursively split the data matrix. It splits columns in pairs by a greedy **G Test** based procedure with threshold ρ : $G(X_i, X_j) = 2 \sum_{x_i \sim X_i} \sum_{x_j \sim X_j} c(x_i, x_j) \cdot \log \frac{c(x_i, x_j) \cdot |T|}{c(x_i) c(x_j)}$ (Figure 1.c); it clusters instances in $|C|$ sets with **online Hard-EM** (Figure 1.b) with cluster number penalty λ : $Pr(\mathbf{X}) = \sum_{C_i \in C} \prod_{X_j \in X} Pr(X_j, C_i)$. Weights are the cluster proportions.



If there are less than m instances, it puts a **naive factorization** over leaves (Figure 1.d). For each univariate distribution it gets its **ML estimation** smoothed by α . LearnSPN hyperparameter space is thus: $\{\rho, \lambda, m, \alpha\}$.

The state-of-the-art, in terms of test likelihood, is **ID-SPN**: it turns LearnSPN in log-likelihood guided expansion of sub-networks approximated by Arithmetic Circuits [8]. However it is overparametrized, and slower.

Tractability is guaranteed if the network size is polynomial in # vars. **Structure quality matters** as much as likelihood. comparing network sizes is more solid than comparing inference times.

LearnSPN is too greedy and the resulting SPNs are overcomplex networks that may not generalize well. **Structure quality desiderata**: smaller but accurate, deeper but not wider, SPNs.

Simplifying by limiting node splits

LearnSPN performs two interleaved **greedy hierarchical** divisive **clustering** processes. Each process benefits from the other one improvements and similarly suffers from the other's mistakes.

Idea: slowing down the processes by limiting the number of nodes to split into. SPN-B, variant of LearnSPN that uses EM for mixture modeling but doing only binary splits for sum nodes children ($k = 2$) when clustering rows.

Objectives: not committing to complex structures too early while retaining same expressive power, indeed successive row splits can represent sum nodes with more than two children; moreover, reducing the node out fan increases the network depth. Plus, there is no need for λ anymore.

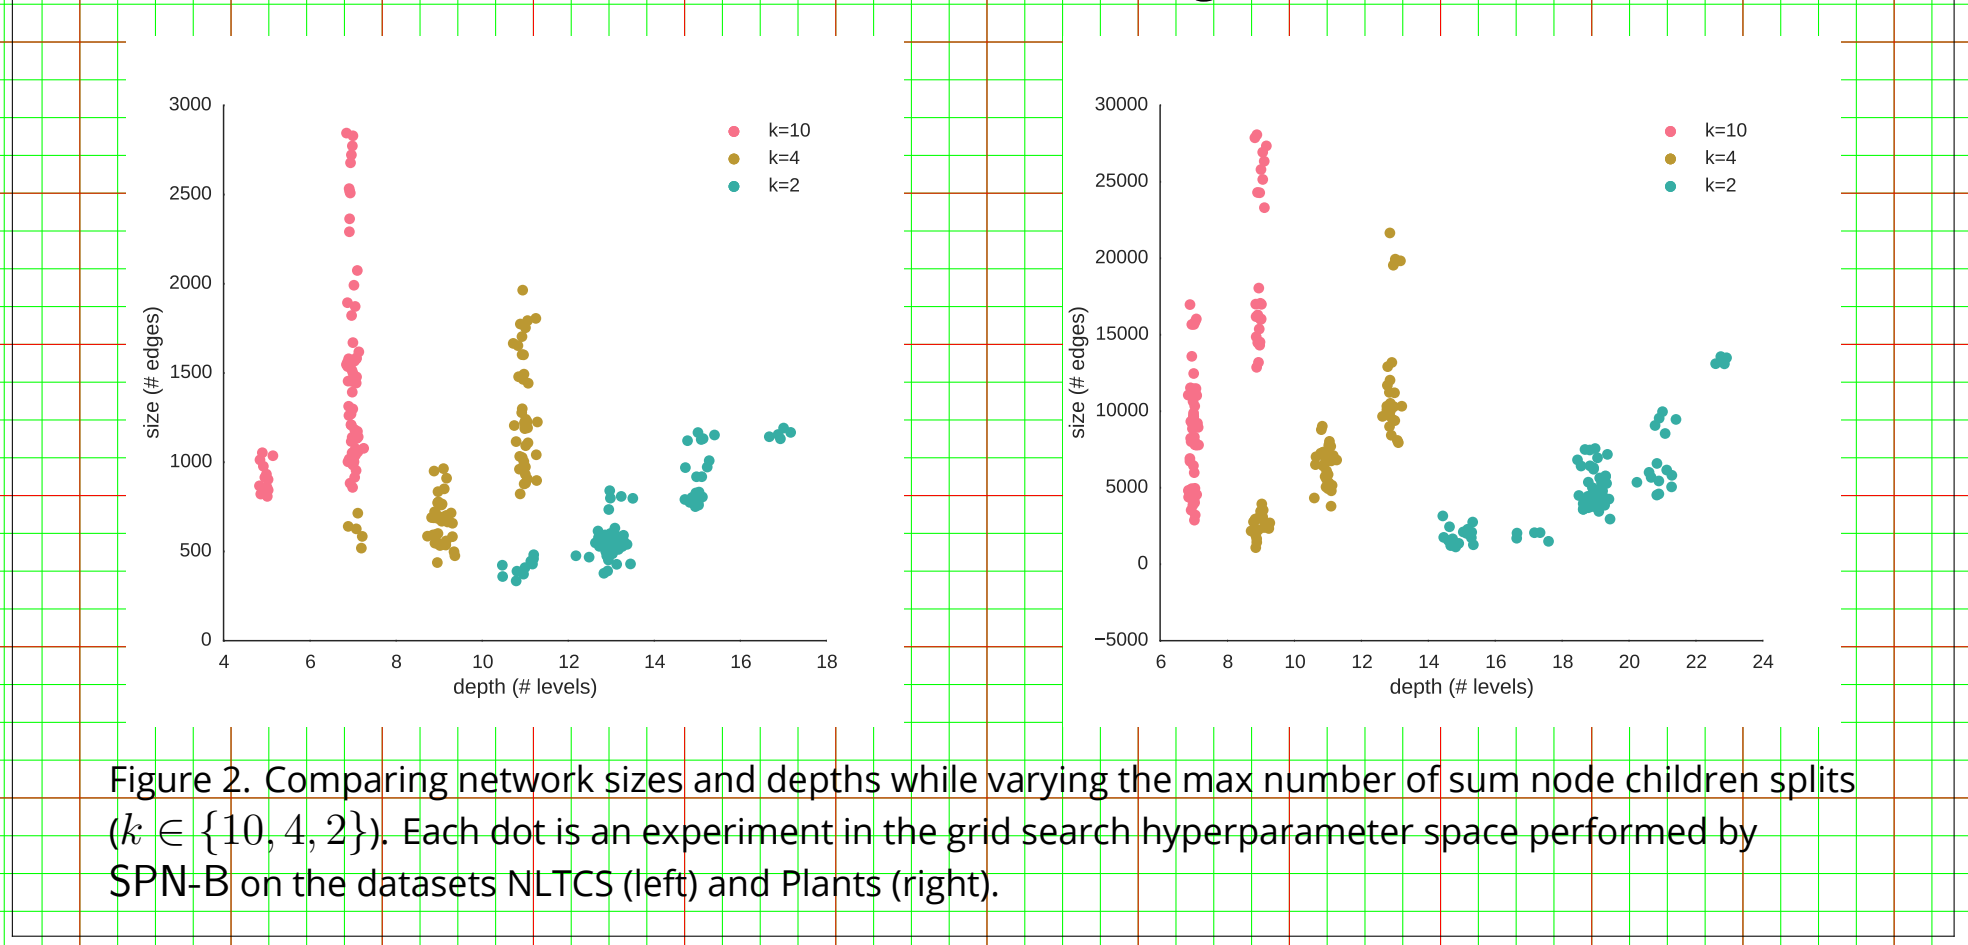
Regularizing by introducing tree distributions as leaves

LearnSPN regularization is is governed by the hyperparameters α and m , however using naive factorizations can be ineffective. In order to get accurate networks, the algorithm prefers smaller values for m , resulting in more complex networks

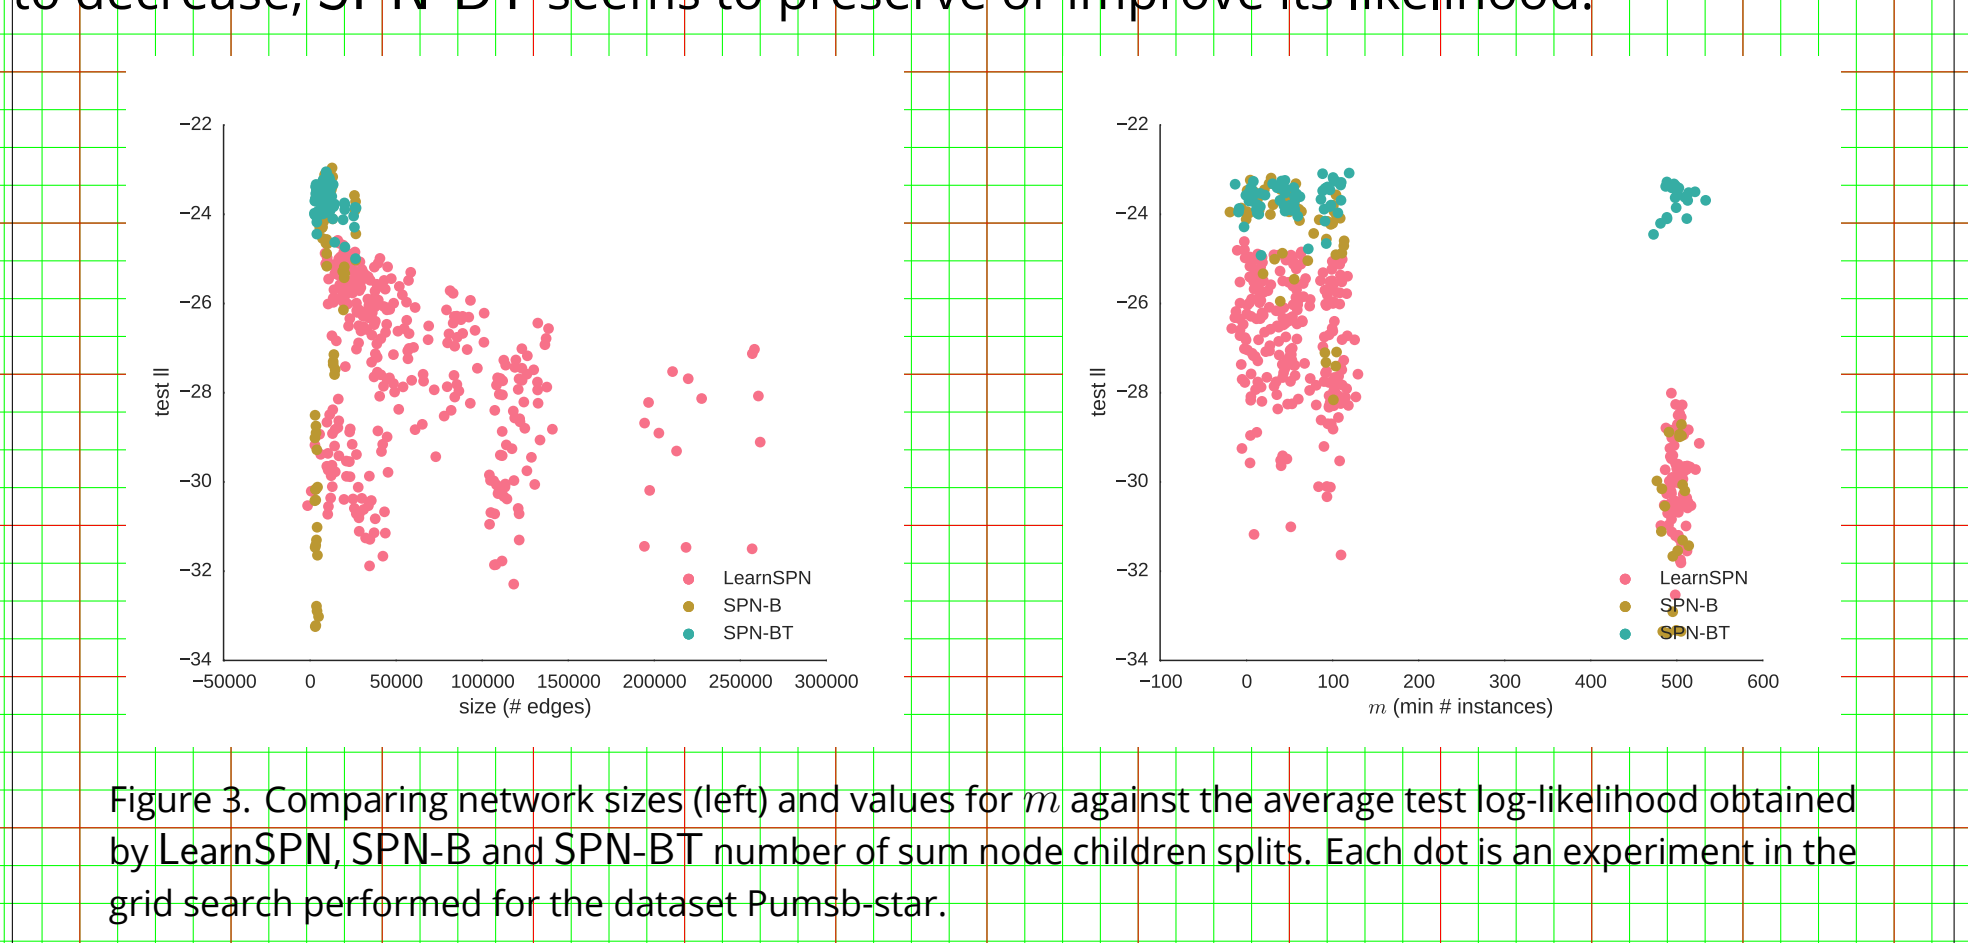
Idea: substitute naive factorizations with Bayesian trees as **multivariate tractable tree distributions**. SPN-BT learns such Trees with the Chow-Liu algorithm while stopping the search.

Objectives: represent more information allowing for larger values of m to be chosen, while preserving tractability for marginals, conditionals and MPE inference.

By increasingly limiting the max number of allowed splits the depth of the structures increases and the network size rate of growth decreases.



SPN-BT reduces the size of the networks even more while preserving SPN-B accuracy. At larger values of m , when both SPN-B and LearnSPN accuracies tend to decrease, SPN-BT seems to preserve or improve its likelihood.



Strengthening by model averaging

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam lobortis facilisis sem. Nullam nec mi et neque pharetra sollicitudin. Praesent imperdiet mi nec ante. Donec ullamcorper, felis non sodales commodo, lectus velit ultrices augue, a dignissim nibh lectus placerat pede. Vivamus nunc nunc, molestie ut, ultricies vel, semper in, velit. Ut porttitor. Praesent in sapien. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Duis fringilla tristique neque. Sed interdum libero ut metus. Pellentesque placerat. Nam rutrum augue a leo. Morbi sed elit sit amet ante lobortis sollicitudin. Praesent blandit blandit mauris. Praesent lectus tellus, aliquet aliquam, luctus a, egestas a, turpis. Mauris lacinia lorem sit amet ipsum. Nunc quis urna dictum turpis accumsan semper.

Experiments

Classical setting for **generative** graphical models structure learning [2]:

- 19 binary datasets from classification, recommendation, frequent pattern mining...[4] [3]
- Training 75% Validation 10% Test 15% splits (no cv)

Comparing both accuracy and structure quality:

- average log-likelihood** on predicting test instances
- networks sizes (# edges)
- network depth (# alternated type layers)

Comparing the state-of-the-art, LearnSPN, ID-SPN and MT [6], against our variations:

- SPN-B using only Binary splits
- SPN-BT with Binary splits and Trees as leaves
- SPN-BB combining Binary splits and Bagging
- SPN-BTB including all variants

Model selection via *grid search* in the same parameter space:

- $\lambda \in \{0.2, 0.4, 0.6, 0.8\}$, $m \in \{1, 50, 100, 500\}$,
- $\rho \in \{5, 10, 15, 20\}$, $\alpha \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$.




Table : Average test log likelihoods for all algorithms. In bold the best values after a Wilcoxon signed rank test with p -value of 0.05.

References

Aaron Dennis and Dan Ventura. "Learning the Architecture of Sum-Product Networks Using Clustering on Variables". In: *Advances in Neural Information Processing Systems* 25. Curran Associates, Inc., 2012, pp. 2033-2041.

Robert Gens and Pedro Domingos. "Learning the Structure of Sum-Product Networks". In: *Proceedings of the 30th International Conference on Machine Learning, JMLR Workshop and Conference Proceedings*, 2013, pp. 873-880.

Jan Van Haaren and Jesse Davis. "Markov Network Structure Learning: A Randomized Feature Generation Approach". In: *Proceedings of the 26th Conference on Artificial Intelligence*. AAAI Press, 2012.

Daniel Lowd and Jesse Davis. "Learning Markov Network Structure with Decision Trees". In: *Proceedings of the 10th IEEE International Conference on Data Mining*. IEEE Computer Society Press, 2010, pp. 334-343.

James Martens and Venkatesh Medaballimi. "On the Expressive Efficiency of Sum Product Networks". In: *CoRR abs/1411.7717* (2014).

Marina Meilă and Michael I. Jordan. "Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1 (2000), pp. 1-48.

Robert Peharz, Bernhard Gelger, and Franz Pernkopf. "Greedy Part-Wise Learning of Sum-Product Networks". In: *Machine Learning and Knowledge Discovery in Databases*. Vol. 8189. LNCS. Springer, 2013, pp. 612-627.

Amirmohammad Rooshenas and Daniel Lowd. "Learning Sum-Product Networks with Direct and Indirect Variable Interactions". In: *Proceedings of the 31st International Conference on Machine Learning, JMLR Workshop and Conference Proceedings*, 2014, pp. 710-718.

ECML-PKDD 2015 - 8th September 2015, Porto, Portugal